

1 Deep-Learning Based Cortical Surface Reconstruction:
2 A Review of Methods, Comparisons, and
3 Recommendations

4 William Ashbee^a, Linlin Lu^a, Harshvardhan Gazula^b, Noah Lewis^a, Satrajit
5 Ghosh^b, Vince D Calhoun^a, Sergey Plis^a

6 ^a*TReNDS Center, Georgia State, Georgia Tech, Emory, Atlanta, GA, USA*

7 ^b*MIT, Cambridge, MA, USA*

8 **Abstract**

Accurate reconstruction of the brain's cortical surface is essential for neuroimaging analyses, enabling precise measurement of cortical structures and supporting studies of neurological conditions. Recent advances in deep learning have significantly improved the accuracy and efficiency of cortical surface reconstruction, offering promising alternatives to traditional methods. In this review, we analyze current deep learning approaches for cortical surface reconstruction, focusing on their design elements and evaluating their performance in terms of accuracy, surface quality, and computational efficiency. Our analysis reveals that some methods achieve high geometric accuracy at the expense of increased computational resources, while others maintain structural integrity with lower resource demands but may sacrifice some accuracy. By systematically evaluating these trade-offs, we identify methods that effectively balance accuracy, structural integrity, and resource efficiency. Our findings indicate that mesh-based methods with optimized loss functions offer the best compromise for most applications. This paper not only benchmarks recent deep learning methods for cortical surface reconstruction but also uncovers underlying design principles that influence performance. We provide actionable guidance on selecting appropriate approaches based on specific needs, highlighting the importance of considering both precision and computational practicality. By emphasizing the interplay between accuracy, resource consumption, and surface quality, we

*Corresponding author

Email addresses: wsashbee@gmail.com (William Ashbee), pku600gt@gmail.com (Linlin Lu), hgazula@mit.edu (Harshvardhan Gazula), lhd231@gmail.com (Noah Lewis), satra@mit.edu (Satrajit Ghosh), vcalhoun@gsu.edu (Vince D Calhoun), s.m.plis@gmail.com (Sergey Plis)

provide a valuable resource for researchers and clinicians in neuroimaging, neuroscience, and related fields.

⁹ *Keywords:* Deep Learning, Cortical Surface Reconstruction, Benchmarking,
¹⁰ Topology, Neuroanatomy, sMRI

¹¹ **1. Introduction**

¹² Understanding the intricate structures of the human brain is fundamental
¹³ to advancing neuroscience and developing treatments for various neurological
¹⁴ disorders. Cortical Surface Reconstruction (CSR) from structural magnetic res-
¹⁵ onance imaging (sMRI) plays a pivotal role in this endeavor by enabling detailed
¹⁶ analysis of cortical morphology, including measurements of cortical thickness,
¹⁷ surface area, and folding patterns [28, 36]. Accurate reconstruction of cortical
¹⁸ surfaces facilitates studies on brain development, neurodegenerative diseases,
¹⁹ and cognitive functions, providing critical insights into the brain’s architecture
²⁰ and its variations across individuals [14, 31].

²¹ FreeSurfer [13] has long been considered the gold standard for CSR, provid-
²² ing reliable reconstructions that serve as ground truth for training deep learning
²³ models [26, 20, 34, 23, 27, 6, 10]. However, *FreeSurfer* is burdened by extended
²⁴ processing times and topological errors, such as self-intersections in the recon-
²⁵ structed surfaces [20, 18]. These limitations have spurred the development of
²⁶ new deep learning-based methods that aim to overcome these challenges.

²⁷ Recent advances in deep learning have ushered in a variety of models specif-
²⁸ ically designed to improve both the efficiency and accuracy of cortical surface
²⁹ generation [26, 20, 34, 23, 27, 6, 10]. The landscape of deep learning-based
³⁰ CSR models can be broadly categorized into three approaches [6]: voxel-based,
³¹ implicit surface-based, and mesh-based methods. Voxel-based approaches typi-
³² cally involve segmenting the sMRI data followed by surface extraction [18, 16] to
³³ generate cortical surfaces. Implicit surface-based frameworks predict occupancy
³⁴ fields or signed distance function (SDF), which guide the surface extraction [10].
³⁵ In contrast, mesh-based methods directly deform a template mesh to align with

36 the sMRI data, leveraging advanced techniques in geometric deep learning to
37 achieve precise reconstructions [20, 23, 34, 6, 27, 26].

38 However, given the diversity of these approaches, there is no clear consensus
39 on which model is best suited for specific tasks or datasets. Moreover, the
40 architectural relationships and functional unit interactions within these models
41 are underexplored, limiting the ability of researchers to make informed decisions
42 when selecting models for their work. To address this gap, we provide
43 a systematic analysis of recent deep learning models for CSR, examining their
44 architectural differences, performance characteristics, and suitability for various
45 applications.

46 This paper makes several key contributions. We evaluate the performance
47 of recent deep learning models for CSR using various metrics such as distance
48 to ground truth, self-intersections, inter-mesh collisions, and overall mesh quality,
49 offering insights into their geometric and topological strengths and weaknesses.
50 We also assess the computational performance of these models, including system
51 memory requirements, processing time, and graphics processing unit (GPU) memory
52 usage, providing practical information for selecting models based on available
53 computational resources. Additionally, we conduct an in-depth examination of
54 these models detailing their architectural choices and functional components
55 including surface extraction methods, mesh deformation techniques, smoothing
56 algorithms, loss functions, topology correction mechanisms, and neural network
57 architectures. By elucidating the architectural relationships, performance
58 characteristics, and key functional units of these models, this work serves as a
59 valuable resource for researchers and clinicians aiming to
60 leverage deep learning for cortical surface reconstruction.

61 The remainder of this paper is organized as follows. In Section 2, we define
62 the core problem of cortical surface reconstruction and review representative
63 deep learning-based approaches for CSR. Section 3 describes the methodology
64 used for our comparative analysis. Section 4 presents the results of our eval-
65 uation, including mesh quality and computational performance measures. In
66 Section 5, we discuss the functional units of CSR frameworks and provide guid-

67 ance on model selection. Finally, Section 6 concludes the paper with a summary
68 of our findings and suggestions for future work.

69 **2. CSR: Problem Statement and Representative Approaches**

70 *2.1. Problem Statement*

71 Cortical surface reconstruction from brain imaging predominantly employs
72 structural, T1-weighted sMRI images to generate surfaces, represented as tri-
73 angular meshes, reflecting the boundaries between white and gray matter, and
74 gray matter and cerebrospinal fluid (CSF), called the white and the pial surface,
75 respectively [13, 28]. The goal is to ensure that these surfaces do not intersect,
76 have no self-intersections, and form contiguous surfaces, while capturing the
77 underlying structure of the brain.

78 Given a three-dimensional structural sMRI image denoted by $\mathbf{I} \in \mathbb{R}^{H \times W \times D}$,
79 we aim to define a transformation ψ that yields a triangular mesh \mathcal{M} :

$$\psi : \mathbf{I} \rightarrow \mathcal{M}, \quad (1)$$

80 where H , W , and D correspond to the height, width, and depth of the image,
81 respectively.

82 The triangular mesh \mathcal{M} is described as a set of vertices \mathbf{V} , edges \mathbf{E} , and
83 faces \mathbf{F} , expressed as $\mathcal{M} = \{\mathbf{V}, \mathbf{E}, \mathbf{F}\}$. Each edge $e \in \mathbf{E}$ is a line segment
84 characterized by its endpoints $\mathbf{u}, \mathbf{v} \in \mathbf{V}$. Similarly, each face $f \in \mathbf{F}$ is defined
85 by an ordered triplet $\mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathbf{V}$, forming a triangular face that is one of the
86 basic units of the mesh surface.

87 The transformation ψ is achieved by first preprocessing an image using a
88 function ρ^* (where $*$ indicates the function is optional or can be the identity
89 transformation), then processing it through a neural network ω parameterized
90 by θ . Finally, the neural network's output is post-processed with a function
91 Δ^* to produce the final mesh. Typically, the function Δ^* involves Laplacian
92 smoothing, although other post-processing functions can also be applied. Thus,

93 ψ is a composition of these three transformations, which can be expressed as:

$$\psi(\mathbf{I}; \boldsymbol{\theta}) = \Delta^*(\omega(\rho^*(\mathbf{I}); \boldsymbol{\theta})) \quad (2)$$

94 Within this framework, researchers train the neural network ω to ascertain
95 the relationship between the preprocessed 3D image and the mesh. Alterna-
96 tively, ψ could encapsulate the operation of the *FreeSurfer* pipeline [13], which
97 the deep learning frameworks are attempting to learn and replace.

98 *2.2. Representative Approaches*

99 In this section, we review representative deep learning-based approaches for
100 CSR providing an overview of their architectural foundations. We review several
101 seminal models that execute modules similar to the *FreeSurfer* CSR pipeline in
102 a fraction of the time.

103 *2.2.1. DeepCSR*

104 *DeepCSR*¹ [10] is an implicit surface framework that uses topology correction
105 and marching cubes [25] for cortical surface generation. It was one of the first
106 deep-learning-based methods for this task. The pipeline involves registering
107 input sMRI and ground truth surfaces to an MNI152 T1 1mm template, using
108 the sMRI as input to predict an implicit surface. After topology correction,
109 marching cubes are applied to preserve topology.

110 The implicit surface \mathcal{S} is defined as:

$$\mathcal{S} = \{\mathbf{p} \in \mathbb{R}^3 \mid d_{\mathcal{S}}(\mathbf{p}) = \ell\}. \quad (3)$$

111 where $d_{\mathcal{S}}(\mathbf{p})$ is the signed distance function (SDF), indicating whether a point
112 is inside or outside \mathcal{S} .

113 The use of implicit functions for surface representation in DeepCSR aligns
114 with the approach pioneered by DeepSDF [29], which demonstrated the effec-
115 tiveness of modeling 3D surfaces as continuous signed distance fields to capture

1Cortical surface reconstruction via deep learning

116 fine geometric details and overcome the limitations of discrete voxel-based meth-
117 ods.

118 A neural network ω predicts $d_S(\mathbf{p})$ using points $\mathcal{P}_{\text{sMRI}}$ sampled from the
119 sMRI space during training. Inference involves uniform sampling within a fixed
120 bounding box, a simpler approach compared to training.

121 For preprocessing, occupancies and distances are computed for $\mathcal{P}_{\text{sMRI}}$ using
122 occupancy grids and the SDF. The grids help identify surface boundaries for
123 marching cubes [10, 25], and distances are calculated with nearest-neighbor
124 searches [4].

125 The architecture consists of an encoder that extracts features from the sMRI
126 using convolutions and pooling, while a hypercolumn module interpolates ac-
127 tivations to create dense feature representations. A decoder then uses these
128 features to predict the implicit surface, employing conditional batch normaliza-
129 tion (CBN) and fully connected layers.

130 Topology correction ensures the implicit surface is genus 0 [2] before march-
131 ing cubes extracts the surface mesh. The framework uses binary cross-entropy
132 (BCE) loss for occupancy fields and L_1 loss for SDF predictions.

133 2.2.2. *TopoFit*

134 *TopoFit* [20] is designed to reduce processing time and computational de-
135 mands by integrating image and graph convolutions to deform a genus 0 mesh
136 template into a subject-specific white matter surface. This approach combines
137 the power of Graph Neural Networks (GNN) [35] and the UNet architecture [33]
138 for efficient mesh adaptation.

139 *TopoFit* uses an image-based UNet to process sMRI data, which is then
140 aligned with mesh coordinates using a Talairach transformation. This alignment
141 enables the creation of image embeddings that capture white matter structure,
142 which are passed to the graph-based UNet. The graph UNet operates on mesh
143 vertex features (coordinates and normals) and predicts deformation vectors to
144 adapt the mesh to the subject’s surface.

145 The deformation process begins with a low-resolution icosphere template,

146 which is refined through seven iterative modules to achieve the final deformed
147 mesh thereby ensuring accurate mesh adaptation to the complex geometry of
148 the white matter surface.

149 To maintain smoothness and geometric consistency, *TopoFit* uses a hinge
150 spring regularization term \mathcal{L}_{reg} on adjacent face normals:

$$\mathcal{L}_{\text{reg}}(\widehat{\mathcal{M}}) = \frac{1}{|\mathbf{E}|} \sum_{(f_a, f_b) \in \mathbf{F}_{\text{adj}}} (1 - \widehat{\mathbf{u}}_{f_a}^T \widehat{\mathbf{v}}_{f_b})^2 \quad (4)$$

151 where $\widehat{\mathbf{u}}_{f_a}$ and $\widehat{\mathbf{v}}_{f_b}$ are the normals of adjacent faces. This regularization pro-
152 motes smooth mesh deformation.

153 Additionally, *TopoFit* uses a localized Chamfer distance $\mathcal{L}_{\text{dist}}$ to measure
154 mesh similarity:

$$\mathcal{L}_{\text{dist}}(\mathcal{M}, \widehat{\mathcal{M}}) = \frac{1}{2|\mathbf{V}|} \sum_{\mathbf{v} \in \mathbf{V}} \min_{n \in \mathcal{N}_{\mathbf{v}}} \|\mathbf{v} - \widehat{\mathbf{v}}_n\| + \|\widehat{\mathbf{v}} - \mathbf{v}_n\|, \quad (5)$$

155 which quantifies the average pairwise distance between mesh vertices. By com-
156 bining $\mathcal{L}_{\text{dist}}$ and \mathcal{L}_{reg} , *TopoFit* ensures smooth, accurate mesh reconstructions
157 with minimal self-intersections.

158 2.2.3. *CorticalFlow*

159 *CorticalFlow* [23] is a neural ODE-based approach for cortical surface recon-
160 struction. It uses a 3D UNet [8] to predict a diffeomorphic flow field, which is
161 then integrated numerically to deform a template mesh. This process ensures
162 a smooth, homeomorphic transformation, preserving topology during deforma-
163 tion.

164 *CorticalFlow* leverages the preprocessing pipeline of *DeepCSR*, training in
165 the MNI152 T1 1mm template space. This choice likely facilitates a fair com-
166 parison between the two frameworks. At inference, only minimal processing is
167 required for the sMRI, as the ground truth surfaces from *FreeSurfer* are not
168 needed.

169 The Diffeomorphic Mesh Deformer (DMD) in *CorticalFlow* uses Euler's
170 ODE solver to integrate through the flow field, ensuring that the mesh re-
171 mains topologically valid. The deformation is governed by the evolution of a

172 time-dependent coordinate map ϕ , described by the first-order ODE:

$$\frac{d\phi(t; \mathbf{V})}{dt} = \psi(\phi(t; \mathbf{V}), t), \text{ with } \phi(0; \mathbf{V}) = \mathbf{V}_0, \quad (6)$$

173 where \mathbf{V}_0 is the initial template mesh configuration, and ψ is the predicted
174 vector field.

175 Training involves three UNets, with each model learning progressively finer
176 deformations. During training, the models are updated sequentially: first train-
177 ing the coarsest model, freezing its weights, then training the next, and so on.
178 The UNets are applied in sequence during inference, repeating three times.

179 The training loss includes Chamfer distance and edge length terms to ensure
180 both accurate vertex positions and smooth mesh edges. The edge length loss
181 for a mesh with vertices \mathbf{V} and edges \mathbf{E} is computed as:

$$L_{\text{edge}} = \frac{1}{|\mathbf{E}|} \sum_{e \in \mathbf{E}} (\|\mathbf{v}_i - \mathbf{v}_j\|_2 - l_{\text{target}})^2 \quad (7)$$

182 *CorticalFlow* guarantees that the deformation remains diffeomorphic and in-
183 vertible, preserving topology, provided the space is Lipschitz continuous and the
184 numerical integration step size is sufficiently small. This ensures no topological
185 defects during deformation [34, 26], though self-intersections can still occur, as
186 noted by Zhao et al. [40], highlighting an area for further exploration.

187 2.2.4. *CorticalFlow++*

188 *CorticalFlow++* [34] builds upon the *CorticalFlow* framework with three key
189 improvements that enhance its robustness:

- 190 1. **ODE Solver Update:** The Runge-Kutta method replaces Euler's method
191 for numerical integration, improving the stability and accuracy of the de-
192 formation.
- 193 2. **Template Mesh Refinement:** The initial template mesh is smoothed
194 and optimized to ensure it includes all vertices from the training set, with
195 no topological defects.
- 196 3. **White-to-Pial Surface Deformation:** *CorticalFlow++* enables the de-
197 formation from the white matter surface to the pial surface, providing
198 explicit correspondences between the two surfaces.

199 In *CorticalFlow++*, correspondence between meshes implies a bijection, where
200 each vertex in one mesh corresponds to a unique vertex in the other. This al-
201 lows for the efficient computation of loss functions such as Mean Squared Error
202 (MSE) or L_1 , both of which have linear complexity $O(N)$, as opposed to Cham-
203 fer Distance’s quadratic complexity $O(N^2)$. This efficiency is particularly useful
204 for large datasets and simplifies the calculation of cortical thickness between the
205 white and pial surfaces.

206 The template creation process in *CorticalFlow++* involves the following
207 steps:

- 208 1. Compute the Signed Distance Function (SDF) for each surface mesh in
209 the training set.
- 210 2. Generate an initial mesh using the marching cubes algorithm on the binary
211 union of these SDFs.
- 212 3. Smooth the mesh with the Laplacian smoothing algorithm and apply De-
213 launay triangulation to obtain a genus-0 template mesh.

214 2.2.5. *Vox2Cortex*

215 *Vox2Cortex* [6] integrates convolutional neural networks (CNNs) and graph
216 neural networks (GNNs) to process sMRI data and generate detailed cortical
217 meshes. Unlike traditional methods that rely on voxel-based reconstruction or
218 implicit surfaces, it avoids techniques like marching cubes and topology correc-
219 tion. The preprocessing pipeline registers images, segmentations, and meshes
220 to MNI152 T1 1mm space, ensuring consistency between the image and the
221 mesh. The iterative mesh refinement strategy used in *Vox2Cortex* is conceptu-
222 ally similar to Pixel2Mesh [38], which deforms an initial 3D mesh into a target
223 shape using a graph-based convolutional approach, leveraging perceptual and
224 geometric loss functions to ensure detailed and smooth surfaces.

225 The framework starts with a UNet for segmenting the sMRI scans to iden-
226 tify cortical structures, and then it deforms a predefined mesh template using
227 cortical features extracted from the scans. The features from each layer of the
228 UNet (across both down- and up-convolution layers) are combined and mapped

229 to mesh vertices, which serve as input to a GNN that predicts the final cortical
230 surface. The GNN updates vertex features by aggregating information from
231 neighboring vertices, following the update rule:

$$\mathbf{h}'_i = \frac{1}{1 + |\mathcal{N}(i)|} \left[\mathbf{W}_0 \mathbf{h}_i + \mathbf{b}_0 + \sum_{j \in \mathcal{N}(i)} (\mathbf{W}_1 \mathbf{h}_j + \mathbf{b}_1) \right] \quad (8)$$

232 where \mathbf{h}'_i is the updated feature for vertex i , and $\mathcal{N}(i)$ represents the neighboring
233 vertices.

234 The loss function in *Vox2Cortex* consists of several components:

235 1. **Segmentation loss \mathcal{L}_{vox} :** Measures the difference between predicted and
236 ground truth segmentation maps, computed using binary cross-entropy
237 (BCE):

$$\mathcal{L}_{\text{vox}}(\widehat{\mathbf{B}}_{\text{seg}}, \mathbf{B}_{\text{seg}}) = \sum_{l=1}^L \mathcal{L}_{\text{BCE}}(\widehat{\mathbf{B}}_l, \mathbf{B}_l), \quad (9)$$

238 2. **Mesh loss $\mathcal{L}_{\text{mesh}}$:** Measures the difference between the predicted and
239 ground truth meshes using consistency terms, including mesh-to-mesh dis-
240 tance, surface normal alignment, and regularization:

$$\mathcal{L}_{\text{mesh}}(\widehat{\mathcal{M}}, \mathcal{M}) = \mathcal{L}_{\text{mesh,cons}}(\widehat{\mathcal{M}}, \mathcal{M}) + \mathcal{L}_{\text{mesh,reg}}(\widehat{\mathcal{M}}), \quad (10)$$

241 where $\mathcal{L}_{\text{mesh,cons}}$ combines curvature-weighted Chamfer distance and inter-
242 mesh normal consistency.

243 3. **Curvature-weighted Chamfer loss \mathcal{L}_C :** Focuses on penalizing discrep-
244 ancies between predicted and true mesh points, with additional weighting
245 based on local curvature to emphasize areas of high detail:

$$\begin{aligned} \mathcal{L}_C(\widehat{\mathcal{M}}_{s,c}, \mathcal{M}_c) &= \frac{1}{|\mathcal{P}_c|} \sum_{\mathbf{u} \in \mathcal{P}_c} \kappa(\mathbf{u}) \min_{\widehat{\mathbf{p}} \in \widehat{\mathcal{P}}_{s,c}} \|\mathbf{u} - \widehat{\mathbf{p}}\|^2 \\ &\quad + \frac{1}{|\widehat{\mathcal{P}}_{s,c}|} \sum_{\widehat{\mathbf{p}} \in \widehat{\mathcal{P}}_{s,c}} \kappa(\tilde{\mathbf{u}}) \min_{\mathbf{u} \in \mathcal{P}_c} \|\widehat{\mathbf{p}} - \mathbf{u}\|^2 \end{aligned} \quad (11)$$

246 where $\kappa(\mathbf{u})$ is the curvature at a given point \mathbf{u} , helping to guide the model's
247 focus toward regions with complex surface features.

248 2.2.6. *CortexODE*

249 *CortexODE* [26] introduces a differential equation to model the evolution
250 of cortical surfaces over time, similar to the *CorticalFlow* formulation of the
251 initial value problem (Equation 6). It uses deformation techniques and provides
252 proofs to minimize topological defects. Notably, *CortexODE* performs topology
253 correction on the signed distance function (SDF) of the initial surface, ensuring
254 the starting surface has genus 0.

255 The *CortexODE* pipeline begins by segmenting the surface into left and right
256 hemisphere white matter regions using a 3D UNet, similar to *Vox2Cortex*. An
257 SDF is computed for the segmentation map, and topology correction is applied
258 to the discrete SDF after Gaussian blurring². Marching cubes are then used to
259 generate an initial surface, which is smoothed using Laplacian smoothing. This
260 initial surface is converted into the white matter surface via the *CortexODE*
261 neural network. The pial surfaces are generated by inflating and smoothing the
262 white matter surface [26].

263 Like *CorticalFlow*, *CortexODE* employs neural ODEs to deform the initial
264 surface during cortical surface reconstruction (CSR). However, unlike *Corti-*
265 *calFlow*, *CortexODE* uses a deformation network that is not a UNet, but instead
266 has a localized receptive field focused on mesh vertices mapped into the sMRI.
267 The deformation network relies on cube sampling, where features are extracted
268 from the sMRI at the mesh vertex locations, reducing the computational cost
269 compared to standard convolutional models like UNet. The sampled cubes are
270 combined into a single matrix of size $|\mathbf{V}| \times C$, where C is the hidden unit size.
271 This matrix is then passed through two fully connected layers to produce a de-
272 formation vector for each mesh vertex. These vectors are used to deform the
273 mesh during the creation of the white and pial surfaces.

274 For training the white matter surface, *CortexODE* uses an initial surface
275 generated by marching cubes and compares it to a ground truth white sur-
276 face from *FreeSurfer*. For training the pial surface, *CortexODE* transforms the

²A custom Python multicore implementation using Numba [21] is provided [26]

277 *FreeSurfer* white surface by inflating and smoothing it, simulating real-world
278 conditions where the white matter surface is available before the pial surface.

279 The inflation and smoothing process for generating the pial surface involves:

280 • **Inflation:** The white matter surface is inflated by adjusting vertex po-
281 sitions along their normal vectors, scaled by a predefined factor. This
282 simulates the anatomical transition from the white matter to the outer
283 cortical boundary (pial surface).

284 • **Smoothing:** Laplacian smoothing is applied to the inflated surface to
285 remove local variations and noise, ensuring a smooth and continuous pial
286 surface.

287 For white surface training, *CortexODE* uses the Chamfer distance loss \mathcal{L}_{ch} :

$$\mathcal{L}_{ch}(\widehat{\mathcal{M}}, \mathcal{M}) = \frac{1}{|\widehat{\mathbf{V}}|} \sum_{\widehat{\mathbf{v}} \in \widehat{\mathbf{V}}} \min_{\mathbf{v} \in \mathbf{V}} \|\widehat{\mathbf{v}} - \mathbf{v}\|^2 + \frac{1}{|\mathbf{V}|} \sum_{\mathbf{v} \in \mathbf{V}} \min_{\widehat{\mathbf{v}} \in \widehat{\mathbf{V}}} \|\mathbf{v} - \widehat{\mathbf{v}}\|^2, \quad (12)$$

288 where $\widehat{\mathcal{M}}$ is the predicted white surface, \mathcal{M} is the ground truth white surface,
289 and $|\widehat{\mathbf{V}}|$ and $|\mathbf{V}|$ are the number of vertices in the predicted and ground truth
290 surfaces, respectively.

291 For pial surface training, *CortexODE* uses mean squared error (MSE) \mathcal{L}_{mse}
292 between the predicted vertex locations $\widehat{\mathbf{V}}$ and the target vertices \mathbf{V} of the ground
293 truth pial surface \mathcal{M} . The use of MSE is feasible because the inflation and
294 smoothing of the *FreeSurfer* white surface ensures a bijective mapping between
295 the white and pial surface vertices.

296 2.2.7. *PialNN*

297 *PialNN* [27] is a partial cortical surface reconstruction framework that takes
298 a white matter surface as input and produces the corresponding pial surface.
299 It generates only one of the two cortical surfaces, requiring the missing sur-
300 face (i.e., white matter) as input. This approach is useful for establishing mesh
301 vertex correspondence between the two surfaces. However, *PialNN* is arguably
302 the most incomplete framework reviewed here because it depends on both an

303 input white matter surface and an sMRI. The input white surface guides the
304 algorithm’s cube sampling, which significantly reduces the number of convolu-
305 tions, making the number of operations proportional to the surface area rather
306 than the sMRI volume. Inspired by approaches such as *Voxel2Mesh* [39], which
307 generate 3D surface meshes directly from volumetric data using an end-to-end
308 trainable architecture and iterative mesh refinement, *PialNN* adopts a similar
309 methodology to accurately construct pial surfaces from MRI volumes.

310 The sMRI volume must be resampled to 256^3 voxels, with 1 mm isotropic
311 resolution and 8-bit depth per voxel ³. For training, *PialNN* requires both the
312 white surface and ground truth pial surface, which are generated by *FreeSurfer*.
313 For inference, *PialNN* requires only the white surface, also generated by *FreeSurfer*.

314 The key components of the *PialNN* architecture are a deformation block and
315 cube sampling. The deformation block iteratively deforms the vertices of the
316 input white mesh to better approximate the pial surface. This is achieved by
317 computing the deformation vectors $\xi\mathbf{V}$ for each vertex and updating the vertex
318 positions as follows:

$$\mathbf{V}_l = \mathbf{V}_{l-1} + \xi\mathbf{V}_{l-1} = \mathbf{V}_{l-1} + \omega(\mathbf{V}_{l-1}, \mathbf{N}_{l-1}, \mathbf{I}), \quad (13)$$

319 where \mathbf{V}_l represents the updated vertex positions. The deformation $\xi\mathbf{V}_{l-1}$ is
320 determined by the deformation network, which uses both localized convolutions
321 and a multi-layer perceptron (MLP) g_{θ_l} . The network processes the previous
322 vertex positions \mathbf{V}_{l-1} , vertex normals \mathbf{N}_{l-1} , and unprocessed sMRI data \mathbf{I} .

323 The network gathers features in two ways: (1) cube sampling to extract
324 local sMRI features and (2) MLP-based point embeddings derived from the
325 mesh vertex coordinates. These point embeddings are generated by processing
326 the vertex coordinates and normals through fully connected layers with leaky
327 ReLU activations. The point embeddings are then combined with image features
328 obtained from cube sampling at multiple scales. These combined features are
329 fed into fully connected layers, which predict the deformation for each vertex

³<https://www.mail-archive.com/freesurfer@nmr.mgh.harvard.edu/msg64274.html>

330 based on its local Cartesian coordinates and the 3D convolutional image features
331 near each vertex.

332 The network uses three deformation blocks in total, followed by Laplacian
333 smoothing to update each vertex, pulling it closer to the weighted average of its
334 neighbors. Finally, the mean squared error (MSE) loss is computed. Since the
335 connectivity of the input and target meshes is identical, MSE is preferred over
336 the more computationally expensive Chamfer distance.

337 *2.2.8. Other Recent Advances in Cortical Surface Reconstruction*

338 During the preparation of this review, several new cortical surface recon-
339 struction frameworks were introduced. While these frameworks are not included
340 in our comparative analysis to maintain focus on the core aspects of CSR, it is
341 important to acknowledge their noteworthy contributions.

342 The authors of *Vox2Cortex*, published two significant modifications to the
343 original framework. In one modification, they developed subject-specific tem-
344 plates that achieve consistent surface correspondence over longitudinal studies
345 involving a single subject [5]. In the other enhancement, a neural ODE to
346 improve the performance of *Vox2Cortex*, analogous to approaches used in *Corti-*
347 *texODE*, *CorticalFlow*, and *CorticalFlow++* was employed [7]. Additionally,
348 *Vox2Cortex* was modified to incorporate an L_1 loss following surface registra-
349 tion for CSR and parcellation, thereby mitigating the limitations associated with
350 Chamfer distance and adding class mappings for generated cortical surfaces [32].

351 To address the limitations of Chamfer distance, Wasserstein distance was
352 employed to generate meshes with more uniformly shaped faces, resulting in
353 improved mesh quality compared to those produced using Chamfer distance
354 alone [22].

355 Ren et al. [30] introduced *FastCSR*, which employs the marching cubes algo-
356 rithm with topology correction similar to that used in *DeepCSR*. Their imple-
357 mentation, which operates in approximately five seconds using a 3D U-Net, is an
358 order of magnitude faster than *DeepCSR*, justifying the nomenclature *FastCSR*.

359 Zheng et al. [42] explore midpoint initial surface deformation between the

360 white and pial surfaces. Furthermore, they also proposed a cycle loss that
361 minimizes the trajectories from the midpoint surfaces to both the white and
362 pial surfaces, ensuring that the Chamfer distance between the midpoint surface
363 and the white and pial surfaces remains equidistant [43].

364 In response to the growing popularity of transformer architectures, An *et*
365 *al.* [1] integrated residual self-attention mechanisms for cortical surface recon-
366 struction. Finally, the domain of cortical surface reconstruction has expanded
367 to include subcortical structures with frameworks such as *Vox2Surf* [19] and
368 *MeshDeform* [41]. These frameworks are capable of constructing surfaces for
369 subcortical structures.

370 3. Methodology

371 By rigorously defining our evaluation criteria and experimental setup, we
372 ensure that our comparison of cortical surface reconstruction methods is both
373 valid and reproducible. This level of methodological transparency is essential for
374 advancing the field, as it allows other researchers to replicate our study and build
375 upon our findings. In this section, we present the methodology used to evaluate
376 the mesh quality and computational performance of the frameworks. Detailed
377 descriptions of the data and code, model training configuration, hardware setup,
378 and containerization and code management are provided in [Appendix A](#).

379 3.1. Mesh Quality Evaluation Methods

380 Evaluating the quality of the reconstructed cortical surfaces is crucial for de-
381 termining the practical applicability of different methods. By employing a range
382 of quantitative metrics including distance measures, self-intersections, and inter-
383 mesh collisions, we provide a thorough assessment that highlights the strengths
384 and limitations of each approach, thereby informing future developments in cor-
385 tical surface reconstruction. These metrics provide an extensive assessment of
386 the structural and topological accuracy of the generated surfaces.

387 3.1.1. *Distance to Ground Truth*

388 Accurately quantifying the deviation of reconstructed cortical surfaces from
389 the ground truth is vital for assessing the performance of different reconstruc-
390 tion methods. By understanding these differences between different methods, we
391 can improve reconstruction algorithms to better capture the true cortical geom-
392 etry, ultimately enhancing the accuracy of neuroscientific research and clinical
393 diagnostics.

394 To assess the accuracy of the predicted meshes, the distances between the
395 predicted meshes and the ground truth meshes were calculated. Since the pre-
396 processing pipelines of the frameworks involved slightly divergent registration
397 techniques, it was necessary to align the meshes to a common space for precise
398 distance calculations. The iterative closest point (ICP) algorithm was employed
399 without necessitating alterations to the frameworks under review, along with
400 scaling and rotation as required for *PialNN*. These procedures yielded a uni-
401 fied linear transformation matrix \mathbf{T} , facilitating the reversion of meshes to their
402 original coordinate system.

403 Conceptually, this process involves four distinct meshes: the *FreeSurfer*
404 ground truth mesh \mathcal{M}_{fs} with vertices \mathbf{V}_{fs} , the framework-specific representa-
405 tion of the ground truth mesh $\mathcal{M}_{\text{framework}}$ with vertices $\mathbf{V}_{\text{framework}}$, the model
406 output mesh $\widehat{\mathcal{M}}_{\text{framework}}$ with vertices $\widehat{\mathbf{V}}_{\text{framework}}$, and the framework-specific
407 mesh prediction registered to *FreeSurfer* space $\widehat{\mathcal{M}}_{\text{fsreg}}$ with vertices $\widehat{\mathbf{V}}_{\text{fsreg}}$.

408 The transformation matrix \mathbf{T} satisfies the equation:

$$\mathbf{V}_{\text{fs}} = \mathbf{T}\mathbf{V}_{\text{framework}} \quad (14)$$

409 Distances between the transformed meshes are computed as:

$$\delta = d_{\text{dist}}(\mathbf{T}\widehat{\mathbf{V}}_{\text{framework}}, \mathbf{T}\mathbf{V}_{\text{framework}}) \quad (15)$$

410 Chamfer distance represents the average value of the closest point distances
411 between two point clouds, while Hausdorff distance denotes the maximum dis-
412 tance of the closest points between two point clouds. Both distances provide a

413 comprehensive view of the mesh comparisons.

414 *3.1.2. Self-Intersections*

415 Self-intersections in the reconstructed meshes can lead to inaccuracies in
416 surface-based analyses and impair the calculation of cortical thickness and other
417 morphological features. By detecting and quantifying self-intersections, we as-
418 sess the topological correctness of each method, which is vital for ensuring the
419 reliability of neuroimaging studies that utilize these surfaces.

420 The assessment of self-intersections employs a KDTree [3] to represent the
421 centers of the mesh triangles. The algorithm iterates through the triangles, and
422 for each center, the k nearest triangles are retrieved using the KDTree to en-
423 hance computational efficiency. Intersections are checked exclusively with these
424 k nearest triangles, and a counter tallies the number of intersections detected.

425 Since self-intersections are invariant with respect to the coordinate system,
426 no coordinate transformations are applied to the predicted meshes prior to cal-
427 culating self-intersections. The medial wall is excluded from topology correction
428 by default in *FreeSurfer*. To ensure a fair comparison, medial wall predictions
429 were excluded from all evaluated models when calculating self-intersections.
430 The exclusion was achieved by utilizing the aparc.a2009s brain atlas and the
431 parcellation scheme proposed by Destrieux et al. [11].

432 *3.1.3. Inter-Mesh Collisions*

433 Inter-mesh collisions between the white and pial surfaces can result in am-
434 biguous or invalid cortical thickness measurements. Evaluating the occurrence
435 of such collisions allows us to determine the suitability of each method for appli-
436 cations that depend on accurate cortical thickness estimation, thereby impacting
437 studies on brain development and disease.

438 Inter-mesh collisions were evaluated to assess the interactions between the
439 white and pial meshes for each hemisphere. Similar to the self-intersection as-
440 sessment, a KDTree was constructed, and the algorithm identified collisions be-
441 tween the white and pial meshes by checking for intersections between triangles.

442 These collisions delineate boundaries where cortical thickness is ambiguously de-
443 fined. The medial wall was excluded here as it was for self-intersections.

444 *3.2. Computational Performance Evaluation Methods*

445 Assessing the computational efficiency of each framework is critical for their
446 practical deployment in research and clinical settings. By measuring system
447 memory usage, execution time, and GPU memory consumption, we provide
448 insights into the resource demands of each method, which can influence their
449 adoption based on available infrastructure.

450 To assess the computational efficiency of each framework, we measured the
451 system memory, execution time, and GPU memory usage. These metrics are es-
452 sential for understanding the resource requirements and scalability of the frame-
453 works under review.

454 *3.2.1. Minimum Required System Memory*

455 Determining the minimum system memory required for each framework helps
456 in understanding their scalability and suitability for different computational
457 environments. This information is particularly important for researchers with
458 limited hardware resources, ensuring that the methods can be effectively utilized
459 across various settings.

460 To establish the minimal system memory requirements for each framework,
461 the lowest stable amount of memory necessary for effective execution was deter-
462 mined. This was achieved by incrementally increasing the allocated memory in
463 steps of 250 MiB, ranging from 250 MiB to 12,750 MiB over 50 runs. Each of the
464 50 runs for a given configuration was tested three times to ensure consistency.

465 *3.2.2. Time Measures*

466 Measuring the execution time of each method provides practical insights
467 into their efficiency and throughput. Faster methods can significantly accelerate
468 neuroimaging pipelines, enabling larger studies and more timely analyses, which
469 is crucial in both research and clinical applications.

470 During the incremental search for minimal memory requirements, the exe-
471 cution time for generating a single surface for one subject was recorded. This
472 measurement reflects the duration from the start to the completion of the pro-
473 gram’s execution. The same subject was used for each run to maintain consis-
474 tency. Each memory allocation configuration was executed three times, totaling
475 150 runs, to mitigate the impact of hardware-induced latency from cold starts.

476 *3.2.3. Peak GPU Memory Usage*

477 Evaluating the peak GPU memory usage of each framework is essential for
478 understanding their compatibility with available hardware, especially for in-
479 stitutions or researchers with varying levels of access to high-end GPUs. This
480 metric can influence the feasibility of adopting certain methods in different com-
481 putational environments.

482 Peak GPU memory usage was monitored during the execution of each frame-
483 work to assess the maximum memory consumption required. This metric is cru-
484 cial for understanding the scalability and resource demands of each framework
485 under evaluation conditions.

486 **4. Results**

487 The results of the mesh quality and computational performance evaluations
488 are presented in this section. All frameworks were evaluated using an identical
489 test set comprising 107 HCP subjects selected from a total of 897 subjects.

490 *4.1. Mesh Quality Measures*

491 The distance plot in Figure 1 provides significant insights into the perfor-
492 mance of various frameworks. Notably, *CorticalFlow++* demonstrates superior
493 performance for pial surfaces. A cluster of highly successful frameworks includes
494 *CorticalFlow*, *CorticalFlow++*, *TopoFit*, and *CortexODE-rk4*. According to the
495 *CortexODE* framework, utilizing the Euler method would require more steps
496 to achieve comparable success to *CortexODE* with the Runge-Kutta 4 (RK4)
497 solver, which is evident in the results [26]. While *DeepCSR* exhibits challenges in

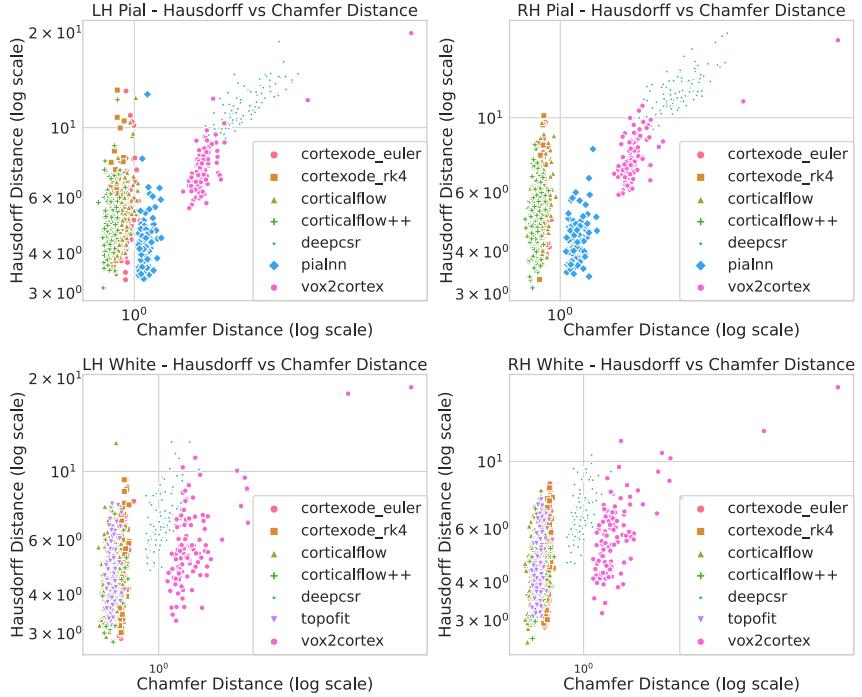


Figure 1: Distribution of distance measures for the test set predicted white and pial surfaces for both hemispheres. Hausdorff distance represents the maximum distance between corresponding points, while Chamfer distance represents the average distance between corresponding points.

generating competitive pial surfaces, it produces nearly competitive white surfaces. *Vox2Cortex* performs suboptimally compared to the state of the art and exhibits catastrophic failures in some instances. *PialNN* generates commendable but not exceptional pial surfaces, with significant improvements observed in the authors' implementation of *CortexODE* using the RK4 solver.

In Figure 2, the component closest point distances are visualized for a single subject. Dark green indicates the best performance, light green the second best, and purple the worst. This figure mirrors the findings in the aggregate measures presented in Figure 1. Increased occurrences of purple and errors are observed in *Vox2Cortex*, *PialNN*, and *DeepCSR*, whereas dark green surfaces are prevalent

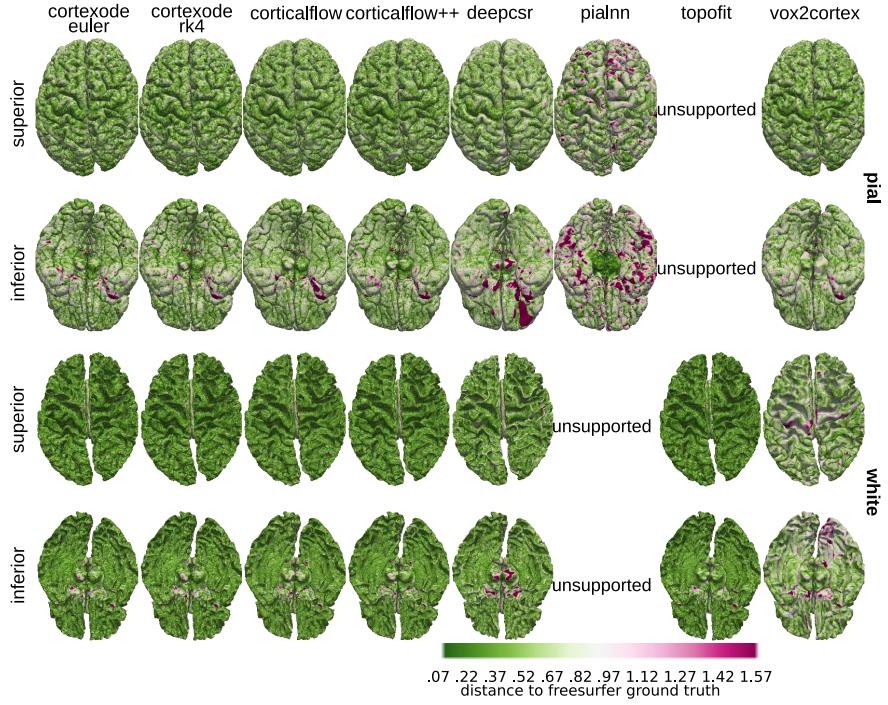


Figure 2: Closest point distances of framework-predicted surfaces to their ground truth for a single subject. Dark green indicates the best performance, light green indicates the next best, and purple indicates the worst performance. *PialNN* and *TopoFit* did not support both surface types at the time of their respective framework publications.

in the cluster of the most successful frameworks—*CortexODE* (Euler and RK4), *TopoFit*, *CorticalFlow*, and *CorticalFlow++*. Overall, the white surfaces exhibit better performance (darker green) than the pial surfaces, which is also reflected in the Chamfer versus Hausdorff plot.

The self-intersection distributions in Figure 3 depict the number of times triangles in the predicted meshes intersect with other triangles across the four surfaces: left and right hemispheres, white and pial. *PialNN* exhibits the highest topological error. While *CorticalFlow* does not perform optimally on pial surface topology given its performance in distance measures, *CorticalFlow++* significantly improves these metrics, potentially due to the use of better templates, the Runge-Kutta 4 solver, or deforming the white surface into a pial surface—likely

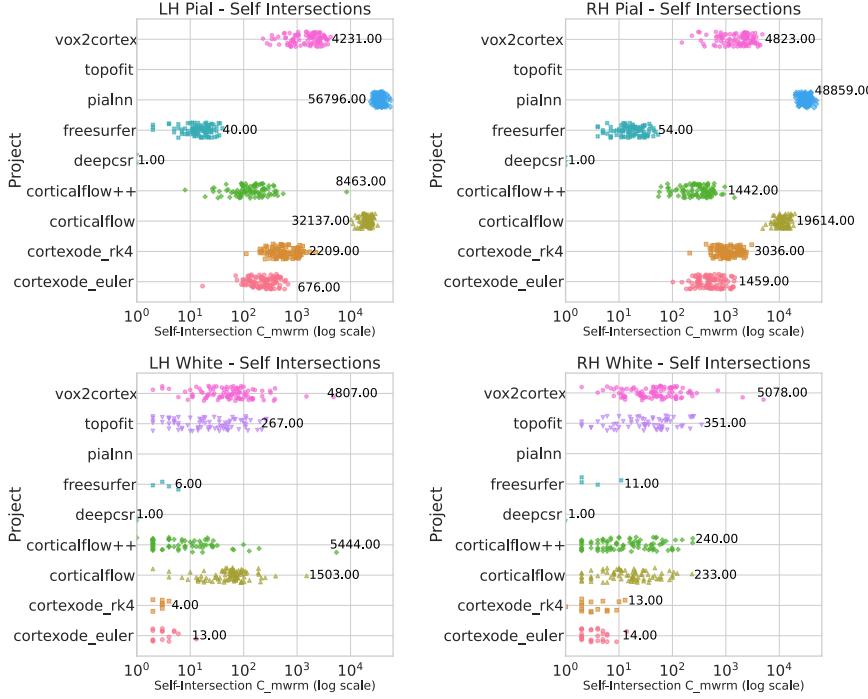


Figure 3: Distribution of collisions and self-intersections for predicted surfaces in the test set (higher values indicate greater severity). These defects can adversely affect geodesic distance measurements within the cortex due to their pervasive distribution. *FreeSurfer* serves as the ground truth, with its measurements provided for reference.

a combination of these factors. White surfaces appear to be generally easier to generate in terms of this measure. *DeepCSR* excels in avoiding self-intersections owing to its topology correction of the SDF and topology-preserving marching cubes algorithm. *CortexODE* produces excellent white surfaces but exhibits moderate performance on pial surfaces in terms of topology, indicating an area for potential improvement in future research. It is noteworthy that *FreeSurfer* also presents some self-intersections. Both *TopoFit* and *Vox2Cortex* perform reasonably well overall on their white surfaces, excluding some extreme outliers observed in *Vox2Cortex*.

It was observed that many defects occur in the sulci and gyri regions where

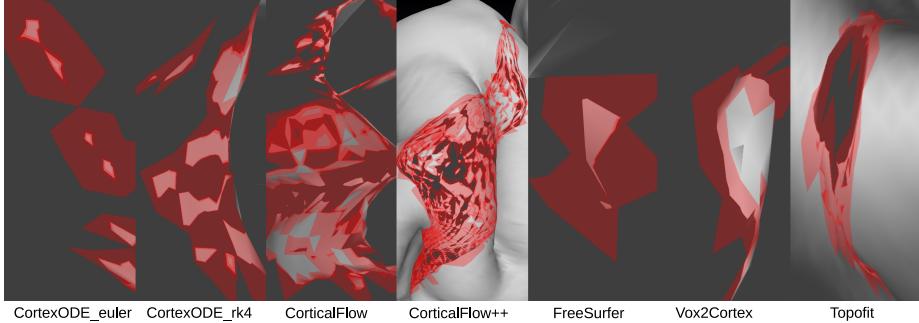


Figure 4: Visualizations and verification of self-intersections in the surfaces of the most affected subjects for selected frameworks. Red highlights indicate areas of self-intersection, light gray represents the outer face, and dark gray represents the inner face. The inversion of outer and inner surface colors within the red boundaries signifies self-intersection.

529 cortical folds arise, as illustrated in some of the failure cases in Figure 4, which
 530 shows self-intersections in red, outward faces in light gray, and inward faces in
 531 dark gray. Typically, defects result from either adjacent folds touching or a
 532 single fold crossing itself. The presence of collisions in a mesh can adversely af-
 533 fect certain types of analyses, such as geodesic distance calculations. Therefore,
 534 assessing the topological quality of these frameworks is crucial for researchers
 535 considering the incorporation of deep learning into their analysis pipelines. The
 536 measures provided in Figure 3 indicate that *DeepCSR* demonstrates exception-
 537 ally high quality in these aspects. This suggests that combining topology cor-
 538 rection with deep learning yields superior results until faster solutions are de-
 539 veloped.

540 For frameworks that produce both surfaces in each hemisphere, an investiga-
 541 tion was conducted to determine whether the white and pial surfaces intersect,
 542 which ideally should not occur. Such white-pial intersections form boundaries
 543 where one surface protrudes into another, resulting in regions where cortical
 544 thickness is poorly defined. The measures in Figure 5 quantify the number of
 545 triangle intersections between the white and pial surfaces in each hemisphere.
 546 It was found that *CortexODE* outperforms other frameworks in these measures,

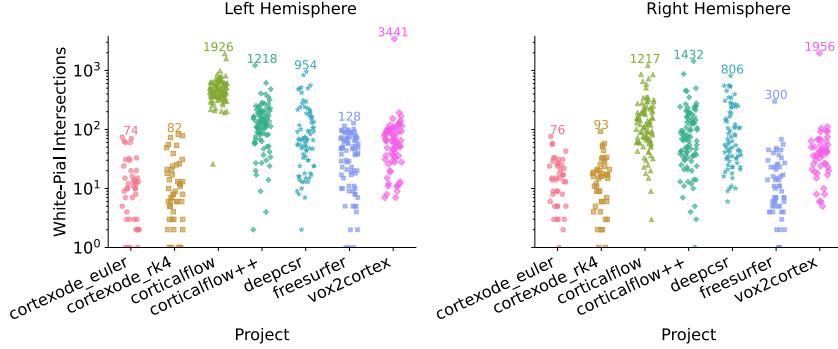


Figure 5: In cortical thickness measurements, white and pial surfaces should not intersect; however, this figure demonstrates that such intersections occasionally occur. The number of triangle-face intersections between these two surfaces is provided for both hemispheres. *FreeSurfer* serves as the ground truth, with its measurements included for reference.

with *FreeSurfer* being the next best. The worst-case outliers depicted in Figure 6 for *CorticalFlow*, *CorticalFlow++*, *DeepCSR*, and *Vox2Cortex* are significant enough to be considered macroscopic features. In the case of *Vox2Cortex*, a catastrophic failure was identified in a surface generated, as shown in Figure 6. Nonetheless, aside from this singular catastrophic failure in each hemisphere, *Vox2Cortex* performs reasonably well in these measures.

4.2. Computational Performance Measures

The computational benchmark results provide insights into the system and GPU memory requirements and the processing time per subject. The results for minimum system memory were generally reasonable across the frameworks, with the exception of *DeepCSR*, as shown in Figure 7A. All frameworks except *DeepCSR* consumed between 2.5 and 5 gigabytes of system memory, which is acceptable for server environments. For client-side deployment, resource availability may vary, and it is advisable to consider both system and GPU memory requirements when determining resource allocation. *DeepCSR* consumed approximately 20 gigabytes of system memory to generate all four surfaces—left and right hemisphere white and pial surfaces. It was observed that providing

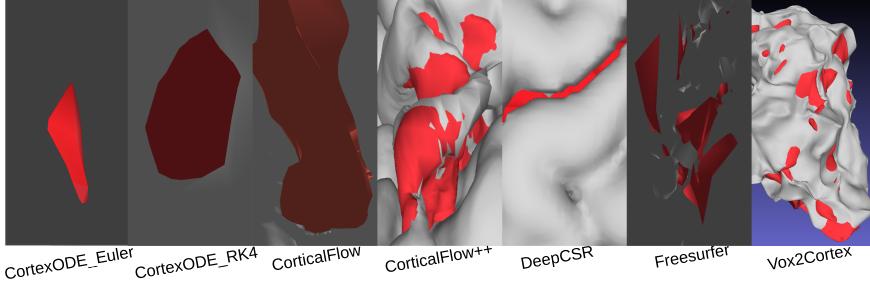


Figure 6: Visualization of interweaving between white and pial surfaces across all frameworks, which should be avoided. One of the two surfaces (white or pial) is colored red to highlight areas where a surface penetrates the other. Defects are located in the lowest quality surfaces for each framework, as indicated in Figure 5. Inter-mesh collisions in the cortex render cortical thickness measurements undefined in these regions (negative thickness values).

slightly less memory than required resulted in significantly slower performance without causing a crash. These findings can assist in deciding appropriate memory allocations for these frameworks. It is important to note that resolution changes can affect these numbers, particularly for *DeepCSR*.

The time measures presented in Figure 7B relate to the generation of cortical surfaces. Notably, *DeepCSR* emerges as the most time-consuming model at a resolution of 256, primarily due to its topology correction processes for implicit surfaces. Additionally, *DeepCSR* employs an iterative loop to sequentially generate sections of the surface during the construction of the implicit surface. Models incorporating ODE solvers, such as *CortexODE* and *CorticalFlow*, are slower at 30 steps compared to *CortexODE* with ten steps, although other architectural factors may also contribute to the performance differences. Their non-ODE counterparts, *PialNN* and *TopoFit*, which involve fewer sequential operations without ODE numerical integration, are slightly faster. Despite these variations, all frameworks achieve the objective of improving upon *FreeSurfer* in terms of speed, reducing processing time from hours to seconds or minutes.

As illustrated in Figure 7C, *DeepCSR* also has the largest GPU memory footprint, despite explicit calls to garbage collection during its iterative prediction of mesh subsections. This is partially attributable to *DeepCSR* generating

583 four surfaces simultaneously. In contrast, *TopoFit*, which generates only a single
 584 surface in this comparison (the left hemisphere white surface), consumes con-
 585 siderable memory (approximately 8 gigabytes). Other frameworks exhibit more
 586 modest memory requirements, with *Vox2Cortex* notably generating all four sur-
 587 faces using approximately 4.5 gigabytes of GPU memory. *PialNN* demonstrates
 588 the lowest memory consumption among the measured frameworks for single sur-
 589 face generation.

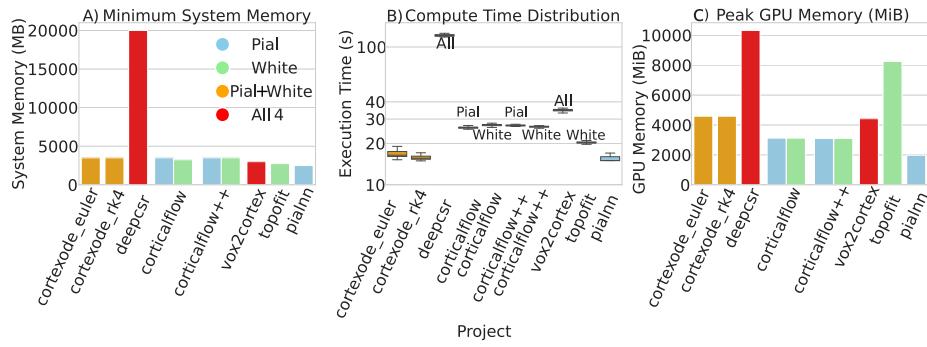


Figure 7: Resource requirements for selected frameworks are depicted in three measures: **A**) minimum system memory, **B**) execution time, and **C**) peak GPU memory usage. The color scheme denotes whether one (white or pial for the left hemisphere), two (white and pial for the left hemisphere), or four (all surfaces for both hemispheres) surfaces were generated. The number of surfaces generated varies based on design choices made by the original authors, which were not altered for this review.

590 The time measures presented in Figure 7B relate to the generation of cor-
 591 tical surfaces. Notably, *DeepCSR* emerges as the most time-consuming model
 592 at a resolution of 256, primarily due to its topology correction processes for im-
 593 plicit surfaces. Additionally, *DeepCSR* employs an iterative loop to sequentially
 594 generate sections of the surface during the construction of the implicit sur-
 595 face. Models incorporating ODE solvers, such as *CortexODE* and *CorticalFlow*,
 596 are slower at 30 steps compared to *CortexODE* with ten steps, although other
 597 architectural factors may also contribute to the performance differences. Their
 598 non-ODE counterparts, *PialNN* and *TopoFit*, which involve fewer sequential op-

599 erations without ODE numerical integration, are slightly faster. Despite these
600 variations, all frameworks achieve the objective of improving upon *FreeSurfer*
601 in terms of speed, reducing processing time from hours to seconds or minutes.

602 It was observed that *DeepCSR* also has the largest GPU memory footprint,
603 despite explicit calls to garbage collection during its iterative prediction of mesh
604 subsections—so much so that For *DeepCSR*, the memory allocation range was
605 extended from 15,000 MiB to 28,500 MiB due to its higher resource demands.
606 This is partially attributable to *DeepCSR* generating four surfaces simulta-
607 neously. In contrast, *TopoFit*, which generates only a single surface in this
608 comparison (the left hemisphere white surface), consumes considerable memory
609 (approximately 8 gigabytes). Other frameworks exhibit more modest memory
610 requirements, with *Vox2Cortex* notably generating all four surfaces using ap-
611 proximately 4.5 gigabytes of GPU memory. *PialNN* demonstrates the lowest
612 memory consumption among the measured frameworks for single surface gener-
613 ation.

614 **5. CSR: A Guide on Functional Units and Their Use**

615 CSR frameworks are composed of various functional units, each responsible
616 for distinct aspects of the reconstruction process. These functional units in-
617 clude surface extraction methods, mesh deformation techniques, smoothing al-
618 gorithms, loss functions, topology correction mechanisms, neural network archi-
619 tectures, and input modalities. Understanding and comparing these functional
620 units across different frameworks is essential for evaluating their strengths, lim-
621 itations, and suitability for specific applications. In this section, we present
622 a comprehensive feature table that categorizes and compares the major CSR
623 frameworks based on their implementation of these functional units, followed
624 by guidance on major considerations for framework selection.

625 *5.1. Functional Units*

626 Our goal is to draw attention to the dominant techniques within the cortical
627 surface reconstruction literature, as depicted in Table 1.

Table 1: Feature Table of Cortical Surface Reconstruction Frameworks

	CortexODE	CorticalFlow	CorticalFlow++	DeepCSR	PialNN	TopoFit	Vox2Cortex
Marching Cubes							
Primary Method				✓			
Initial surface	✓						
Mesh Deformation							
ODE Solver	✓	✓	✓				
Without Solver				✓	✓	✓	
Smoothing							
Initial Surface	✓		✓			✓	
Output Surface	✓			✓			
Loss term					✓	✓	
Loss							
Chamfer distance	✓	✓	✓		✓ _{Local}	✓ _{Curve}	
Other Distance	✓ _{MSE}			✓ _{L₁}	✓ _{MSE}		
Cross Entropy	✓			✓		✓	
Inter/Intra Mesh Normal						✓	
Edge Length	✓	✓	✓			✓	
Laplacian Loss					✓		
Face Normals (Hinge Spring)				✓			
Topology Correction							
Implicit function (not mesh)	✓			✓			
Neural Net Architecture							
Graph Neural Networks					✓	✓	
UNet	✓	✓	✓		✓	✓	
Refinement Networks	✓		✓		✓		
Cube Sampling	✓				✓		
ODE Solver	✓	✓	✓				
Hypercolumn				✓		✓	
Input							
sMRI Only	✓ _{white}	✓	✓ _{white}	✓	✓ _{white}	✓	
sMRI + white surface	✓ _{pial}		✓ _{pial}		✓ _{pial}		

628 **Marching Cubes:** DeepCSR and CortexODE use marching cubes for initial
 629 surface extraction [10, 26]. Vox2Cortex moves away from marching cubes due
 630 to stair-step artifacts [6].

631 **Mesh Deformation:** Deformation frameworks like CortexODE, Corti-
 632 calFlow, and CorticalFlow++ integrate ODE solvers for continuous transfor-
 633 mations [26, 23, 34]. PialNN, TopoFit, and Vox2Cortex employ mesh deformation
 634 without ODE solvers, relying on direct vertex updates [27, 20, 6].

635 **Smoothing:** Smoothing techniques are used to mitigate self-intersections
 636 and ensure smooth surfaces. Frameworks such as PialNN, CortexODE, and

637 *Vox2Cortex* incorporate smoothing at various stages [27, 26, 6].

638 These frameworks have the following smoothing strategies:

639 • *PialNN*: Applied to output surface.

640 • *CortexODE*: Used for the initial and output surfaces.

641 • *CorticalFlow++*: Employed during template creation.

642 • *TopoFit* and *Vox2Cortex*: Incorporated within the loss function.

643 The *TopoFit* and *FreeSurfer* related papers [20, 15, 14] highlight the inadequacy of some smoothing terms in optimization.

644 **Loss Functions:** Loss functions are crucial for training frameworks, with
645 Chamfer distance being prominent in many [12]. *Vox2Cortex* uses a combination
646 of Chamfer distance, edge length, and Laplacian loss [6]. Loss functions should
647 ideally possess the following attributes [12]:

649 • Differentiability with respect to sampled points,

650 • Computational efficiency, and

651 • Robustness against outliers.

652 Chamfer distance is useful over short deformations, but requires additional
653 smoothing terms over larger deformations [20, 26, 23].

654 MSE and L_1 losses provide a stricter structural constraint than chamfer and
655 can be used without smoothing loss terms over short deformations especially if
656 smoothing is applied to the mesh after it is generated [27, 26, 10].

657 Cross Entropy is used for preliminary segmentation tasks.

658 Additional loss terms ensure mesh quality:

659 • *Vox2Cortex* uses intra and inter-mesh vertex normals for mesh deformation
660 consistency [6].

661 • *CorticalFlow* and *Vox2Cortex* use edge length loss to impose structural
662 constraints [23, 6].

663 • *Vox2Cortex* employs Laplacian loss to encourage smooth deformations [6].

664 • *TopoFit* uses a face normal hinge spring term to maintain smooth surface
665 normals [20].

666 **Topology Correction:** Topology correction ensures the precursor implicit
667 surface representations are homeomorphic to a sphere. *DeepCSR* and *CortexODE*
668 perform topology correction on implicit functions, not directly on
669 meshes [10, 26].

670 **Neural Net Architecture:** Architectures include GNNs, UNets, refine-
671 ment networks, cube sampling, and hypercolumns. Frameworks like *Vox2Cortex*
672 and *TopoFit* leverage GNNs for mesh deformation [6, 20]. UNet architectures
673 are employed for feature extraction in frameworks such as *CortexODE* and
674 *TopoFit* [9, 33, 24, 6, 20, 23, 34, 26]. Refinement networks replicate weights
675 to perform the same task multiple times at increasing levels of precision [27,
676 23, 34]. Cube sampling enables local receptive field in the sMRI [27, 26]. Hy-
677 percolumns behave like residual connections to preserve high frequency features
678 [17, 10, 6, 7, 37].

679 **Input:** Some frameworks require only sMRI as input, while others require
680 sMRI and existing white surface meshes. *CorticalFlow++* and *CortexODE* com-
681 bine both approaches for different surface types [34, 26].

682 5.2. Recommendations

683 When choosing a CSR framework, understanding trade-offs between speed,
684 topological accuracy, and loss function complexity is essential. Frameworks
685 like *CorticalFlow++*, *CortexODE*, and *TopoFit* provide fast, high-accuracy re-
686 constructions, suitable for applications needing quick results without sacrific-
687 ing quality. For the highest topological accuracy, *DeepCSR* is recommended,
688 though it requires longer processing times. *Vox2Cortex* explores advanced loss
689 functions, such as curvature-weighted Chamfer distance, but the architecture it-
690 self did not perform well. Simpler frameworks like *PialNN* and *CortexODE* use
691 straightforward loss functions with Laplacian smoothing, maintaining smooth

692 surfaces with less complexity, though only *CortexODE* can independently per-
693 form CSR for both surfaces without *FreeSurfer*.

- 694 • **Computational Resources:** Frameworks like *CorticalFlow++* are ideal
695 for low-memory environments, especially with lower-resolution meshes.
696 For projects with high computational resources, *CortexODE* and *TopoFit*
697 provide finer mesh details and accuracy, while *DeepCSR* remains memory-
698 intensive at high resolutions but is recommended where topological accu-
699 racy is a priority.
- 700 • **Desired Accuracy and Mesh Quality:** For studies demanding pre-
701 cise anatomical fidelity, frameworks like *CortexODE*, *CorticalFlow++*,
702 and *TopoFit* with advanced loss functions and topology correction are
703 ideal. For a balance between speed and accuracy, *CorticalFlow++* pro-
704 vides reliable reconstructions with efficient processing, making it suitable
705 for general-purpose applications.
- 706 • **Model Complexity and Ease of Use:** *PialNN* offers the simplest
707 implementation, though it relies on *FreeSurfer* for initial inputs. *Cor-*
708 *texODE* extends this by enabling independent operation from *FreeSurfer*.
709 Flexible and extensible options, such as *TopoFit*, *CortexODE*, and *Corti-*
710 *calFlow++*, allow customization for specific research needs. *DeepCSR* is
711 well-documented and user-friendly, making deployment straightforward,
712 with other frameworks also compatible with most workflows.

713 **6. Conclusion**

714 This review of deep learning-based CSR has revealed several key insights
715 into the current state and future directions of the field. The emergence of
716 diverse deep learning models has introduced various architectural approaches,
717 yet determining which models are most effective in specific contexts remains
718 an open question. Our analysis underscores the importance of leveraging high-
719 dimensional features to enhance inference quality, as demonstrated by models

720 like *DeepCSR* with its hypercolumn features and the use of skip connections in
721 UNet architectures. Additionally, the implementation of refinement networks
722 across multiple stages, allowing for coarse-to-fine mesh generation, has proven
723 effective in improving surface quality.

724 One of the significant findings from our results is that *CorticalFlow++*
725 emerged as the leading framework for generating pial surfaces, achieving a su-
726 perior balance between accuracy and computational efficiency. *CortexODE_rk4*,
727 *TopoFit*, and *CorticalFlow* also exhibited strong performance, particularly in
728 terms of surface accuracy and quality. However, *DeepCSR*, while showing com-
729 mendable results for white surfaces and excelling in minimizing self-intersections
730 due to its topology correction mechanisms, faced challenges with pial surfaces
731 and required substantially more computational resources. Conversely, *Vox2Cortex*
732 displayed several catastrophic failures, highlighting variability in model robust-
733 ness and the need for further refinement.

734 The practical implications of these findings are significant for both research
735 and clinical applications. By providing a detailed comparison of CSR frame-
736 works, we equip researchers and clinicians with the knowledge to select the
737 most appropriate tools based on specific needs such as computational resources,
738 desired accuracy, and ease of integration into existing workflows. This guidance
739 can streamline the decision-making process, optimize resource allocation, and
740 potentially accelerate the adoption of deep learning methods in neuroimaging
741 studies.

742 Our work contributes to the field by uncovering underlying design principles
743 that influence CSR performance. By emphasizing the importance of architec-
744 tural choices—such as the use of Graph Neural Networks (GNNs), loss function
745 designs, and topology correction mechanisms—we offer actionable insights that
746 can inform the development of next-generation CSR models. These insights can
747 lead to improvements in cortical surface reconstruction accuracy, efficiency, and
748 robustness, ultimately enhancing the quality of neuroimaging analyses.

749 In conclusion, deep learning-based CSR models have demonstrated great po-
750 tential, often rivaling or surpassing traditional methods like *FreeSurfer* in terms

of speed and accuracy. The advances highlighted in this review not only facilitate a better understanding of the brain's architecture but also open avenues for more precise diagnosis and treatment of neurological conditions. By addressing the challenges of balancing precision with computational practicality, this work paves the way for broader adoption and continued innovation in the field of neuroimaging.

References

- [1] An, M., Chen, J., Cao, Y., Tao, K., Wang, J., Wang, C., Zhao, K., Liu, Z., 2023. Resattn-recon: Residual self-attention based cortical surface reconstruction. *Frontiers in Physics* 11, 172.
- [2] Bazin, P.L., Pham, D.L., 2007. Topology correction of segmented medical images using a fast marching algorithm. *Computer methods and programs in biomedicine* 88, 182–190.
- [3] Bentley, J.L., 1975. Multidimensional binary search trees used for associative searching. *Commun. ACM* 18, 509–517. URL: <https://doi.org/10.1145/361002.361007>, doi:[10.1145/361002.361007](https://doi.org/10.1145/361002.361007).
- [4] Bentley, J.L., 1990. K-d trees for semidynamic point sets, in: *Proceedings of the sixth annual symposium on Computational geometry*, pp. 187–197.
- [5] Bongratz, F., Fecht, J., Rickmann, A.M., Wachinger, C., 2024a. V2c-long: Longitudinal cortex reconstruction with spatiotemporal correspondence. *arXiv preprint arXiv:2402.17438* .
- [6] Bongratz, F., Rickmann, A.M., Pölsterl, S., Wachinger, C., 2022. Vox2cortex: fast explicit reconstruction of cortical surfaces from 3d mri scans with geometric deep neural networks, in: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 20773–20783.

- 777 [7] Bongratz, F., Rickmann, A.M., Wachinger, C., 2024b. Neural deformation fields for template-based reconstruction of cortical surfaces from mri.
778 Medical Image Analysis 93, 103093.
- 780 [8] Chen, R.T., Rubanova, Y., Bettencourt, J., Duvenaud, D.K., 2018. Neural
781 ordinary differential equations. Advances in neural information processing
782 systems 31.
- 783 [9] Çiçek, Ö., Abdulkadir, A., Lienkamp, S.S., Brox, T., Ronneberger, O.,
784 2016. 3d u-net: learning dense volumetric segmentation from sparse annotation,
785 in: Medical Image Computing and Computer-Assisted Intervention–
786 MICCAI 2016: 19th International Conference, Athens, Greece, October
787 17-21, 2016, Proceedings, Part II 19, Springer. pp. 424–432.
- 788 [10] Cruz, R.S., Lebrat, L., Bourgeat, P., Fookes, C., Fripp, J., Salvado, O.,
789 2021. DeepCSR: A 3d deep learning approach for cortical surface reconstruc-
790 tion, in: Proceedings of the IEEE/CVF Winter Conference on Applications
791 of Computer Vision, pp. 806–815.
- 792 [11] Destrieux, C., Fischl, B., Dale, A., Halgren, E., 2010. Automatic parcella-
793 tion of human cortical gyri and sulci using standard anatomical nomencla-
794 ture. Neuroimage 53, 1–15.
- 795 [12] Fan, H., Su, H., Guibas, L.J., 2017. A point set generation network for
796 3d object reconstruction from a single image, in: Proceedings of the IEEE
797 conference on computer vision and pattern recognition, pp. 605–613.
- 798 [13] Fischl, B., 2012. Freesurfer. Neuroimage 62, 774–781.
- 799 [14] Fischl, B., Dale, A.M., 2000. Measuring the thickness of the human cere-
800 bral cortex from magnetic resonance images. Proceedings of the National
801 Academy of Sciences 97, 11050–11055.
- 802 [15] Fischl, B., Sereno, M.I., Dale, A.M., 1999. Cortical surface-based analysis:
803 Ii: inflation, flattening, and a surface-based coordinate system. Neuroimage
804 9, 195–207.

- 805 [16] Gopinath, K., Desrosiers, C., Lombaert, H., 2021. Segrecon: Learning
806 joint brain surface reconstruction and segmentation from images, in: Medical
807 Image Computing and Computer Assisted Intervention–MICCAI 2021:
808 24th International Conference, Strasbourg, France, September 27–October
809 1, 2021, Proceedings, Part VII 24, Springer. pp. 650–659.
- 810 [17] Hariharan, B., Arbeláez, P., Girshick, R., Malik, J., 2015. Hypercolumns
811 for object segmentation and fine-grained localization, in: Proceedings of the
812 IEEE conference on computer vision and pattern recognition, pp. 447–456.
- 813 [18] Henschel, L., Conjeti, S., Estrada, S., Diers, K., Fischl, B., Reuter, M.,
814 2020. Fastsurfer – a fast and accurate deep learning based neuroimaging
815 pipeline. NeuroImage 219, 117012.
- 816 [19] Hong, Y., Ahmad, S., Wu, Y., Liu, S., Yap, P.T., 2021. Vox2surf: Implicit
817 surface reconstruction from volumetric data, in: Machine Learning
818 in Medical Imaging: 12th International Workshop, MLMI 2021, Held in
819 Conjunction with MICCAI 2021, Strasbourg, France, September 27, 2021,
820 Proceedings 12, Springer. pp. 644–653.
- 821 [20] Hoopes, A., Iglesias, J.E., Fischl, B., Greve, D., Dalca, A.V., 2022. Topofit:
822 rapid reconstruction of topologically-correct cortical surfaces. Proceedings
823 of machine learning research 172, 508.
- 824 [21] Lam, S.K., Pitrou, A., Seibert, S., 2015. Numba: A llvm-based python jit
825 compiler, in: Proceedings of the Second Workshop on the LLVM Compiler
826 Infrastructure in HPC, pp. 1–6.
- 827 [22] Le, T., Nguyen, K., Sun, S., Han, K., Ho, N., Xie, X., 2023. Diffeomorphic
828 deformation via sliced wasserstein distance optimization for cortical surface
829 reconstruction. arXiv preprint arXiv:2305.17555 .
- 830 [23] Lebrat, L., Santa Cruz, R., de Gournay, F., Fu, D., Bourgeat, P., Fripp, J.,
831 Fookes, C., Salvado, O., 2021. Corticalflow: a diffeomorphic mesh trans-

- 832 former network for cortical surface reconstruction. Advances in Neural
833 Information Processing Systems 34, 29491–29505.
- 834 [24] Liu, Z., Tong, L., Chen, L., Jiang, Z., Zhou, F., Zhang, Q., Zhang, X.,
835 Jin, Y., Zhou, H., 2023. Deep learning based brain tumor segmentation: a
836 survey. Complex & intelligent systems 9, 1001–1026.
- 837 [25] Lorensen, W.E., Cline, H.E., 1987. Marching cubes: A high resolution
838 3d surface construction algorithm. ACM siggraph computer graphics 21,
839 163–169.
- 840 [26] Ma, Q., Li, L., Robinson, E.C., Kainz, B., Rueckert, D., Alansary, A., 2022.
841 Cortexode: Learning cortical surface reconstruction by neural odes. IEEE
842 Transactions on Medical Imaging .
- 843 [27] Ma, Q., Robinson, E.C., Kainz, B., Rueckert, D., Alansary, A., 2021. Pi-
844 alnn: a fast deep learning framework for cortical pial surface reconstruction,
845 in: Machine Learning in Clinical Neuroimaging: 4th International Work-
846 shop, MLCN 2021, Held in Conjunction with MICCAI 2021, Strasbourg,
847 France, September 27, 2021, Proceedings 4, Springer International Pub-
848 lishing. pp. 73–81.
- 849 [28] Makropoulos, A., Robinson, E.C., Schuh, A., Wright, R., Fitzgibbon, S.,
850 Bozek, J., Counsell, S.J., Steinweg, J., Vecchiato, K., Passerat-Palmbach,
851 J., et al., 2018. The developing human connectome project: A minimal
852 processing pipeline for neonatal cortical surface reconstruction. Neuroimage
853 173, 88–112.
- 854 [29] Park, J.J., Florence, P., Straub, J., Newcombe, R., Lovegrove, S., 2019.
855 Deepsdf: Learning continuous signed distance functions for shape represen-
856 tation, in: Proceedings of the IEEE/CVF conference on computer vision
857 and pattern recognition, pp. 165–174.
- 858 [30] Ren, J., Hu, Q., Wang, W., Zhang, W., Hubbard, C.S., Zhang, P., An,

- 859 N., Zhou, Y., Dahmani, L., Wang, D., et al., 2022. Fast cortical surface
860 reconstruction from mri using deep learning. *Brain Informatics* 9, 6.
- 861 [31] Riccelli, R., Toschi, N., Nigro, S., Terracciano, A., Passamonti, L., 2017.
862 Surface-based morphometry reveals the neuroanatomical basis of the five-
863 factor model of personality. *Social cognitive and affective neuroscience* 12,
864 671–684.
- 865 [32] Rickmann, A.M., Bongratz, F., Wachinger, C., 2023. Vertex correspon-
866 dence in cortical surface reconstruction, in: *International Conference on*
867 *Medical Image Computing and Computer-Assisted Intervention*, Springer.
868 pp. 318–327.
- 869 [33] Ronneberger, O., Fischer, P., Brox, T., 2015. U-net: Convolutional net-
870 works for biomedical image segmentation, in: *Medical Image Computing*
871 and *Computer-Assisted Intervention–MICCAI 2015: 18th International*
872 *Conference*, Munich, Germany, October 5–9, 2015, *Proceedings, Part III*
873 18, Springer. pp. 234–241.
- 874 [34] Santa Cruz, R., Lebrat, L., Fu, D., Bourgeat, P., Fripp, J., Fookes, C., Sal-
875 vado, O., 2022. Corticalflow++: Boosting cortical surface reconstruction
876 accuracy, regularity, and interoperability, in: *International Conference on*
877 *Medical Image Computing and Computer-Assisted Intervention*, Springer.
878 pp. 496–505.
- 879 [35] Scarselli, F., Gori, M., Tsai, A.C., Hagenbuchner, M., Monfardini, G., 2008.
880 The graph neural network model. *IEEE transactions on neural networks*
881 20, 61–80.
- 882 [36] Tamnes, C.K., Herting, M.M., Goddings, A.L., Meuwese, R., Blakemore,
883 S.J., Dahl, R.E., Güroğlu, B., Raznahan, A., Sowell, E.R., Crone, E.A.,
884 et al., 2017. Development of the cerebral cortex across adolescence: a
885 multisample study of inter-related longitudinal changes in cortical volume,
886 surface area, and thickness. *Journal of Neuroscience* 37, 3402–3412.

- 887 [37] Ts'o, D.Y., Zarella, M., Burkitt, G., 2009. Whither the hypercolumn? The
888 Journal of physiology 587, 2791–2805.
- 889 [38] Wang, N., Zhang, Y., Li, Z., Fu, Y., Liu, W., Jiang, Y.G., 2018. Pixel2mesh:
890 Generating 3d mesh models from single rgb images, in: Proceedings of the
891 European conference on computer vision (ECCV), pp. 52–67.
- 892 [39] Wickramasinghe, U., Remelli, E., Knott, G., Fua, P., 2020. Voxel2mesh:
893 3d mesh model generation from volumetric data, in: Medical Image Com-
894 puting and Computer Assisted Intervention—MICCAI 2020: 23rd Interna-
895 tional Conference, Lima, Peru, October 4–8, 2020, Proceedings, Part IV
896 23, Springer. pp. 299–308.
- 897 [40] Zhao, F., Wu, Z., Li, G., 2023a. Deep learning in cortical surface-
898 based neuroimage analysis: a systematic review. Intelligent Medicine
899 3, 46–58. URL: <https://www.sciencedirect.com/science/article/pii/S2667102622000493>, doi:<https://doi.org/10.1016/j.imed.2022.06.002>.
- 900 [41] Zhao, J., Liu, S., Ahmad, S., Yap, P.T., 2023b. Meshdeform: Surface recon-
901 struction of subcortical structures via human brain mri, in: International
902 Conference on Information Processing in Medical Imaging, Springer. pp.
903 536–547.
- 904 [42] Zheng, H., Li, H., Fan, Y., 2023. Surfmn: Joint reconstruction of mul-
905 tiple cortical surfaces from magnetic resonance images. arXiv preprint
906 arXiv:2303.02922 .
- 907 [43] Zheng, H., Li, H., Fan, Y., 2024. Coupled reconstruction of cortical sur-
908 faces by diffeomorphic mesh deformation. Advances in Neural Information
909 Processing Systems 36.

912 **Appendix A. Additional Methodology Details**

913 *Appendix A.1. Data and Code*

914 The data utilized in this study were obtained from the Human Connectome
915 Project (HCP). The specific datasets and subjects accessed from the HCP are
916 available through ConnectomeDB at <https://db.humanconnectome.org/>.

917 To ensure reproducibility and facilitate further research, the code developed
918 for the analyses presented in this study is available on GitHub⁴. This reposi-
919 tory includes links to additional repositories where the original frameworks were
920 forked and modified for analysis, benchmarking, and deployment with Singular-
921 ity.

922 For each framework, HCP data processed by *FreeSurfer* were used as the
923 starting point, followed by the application of framework-specific preprocessing
924 scripts, as documented in the respective framework overviews.

925 *Appendix A.2. Model Training and Configuration*

926 *Appendix A.2.1. Model Training*

927 Framework-specific instructions and default configurations, as recommended
928 in their respective README files and configuration files, were adhered to dur-
929 ing model training. Given that architectural choices, schedulers, learning rates,
930 and other design parameters are interdependent, the responsibility for repro-
931 ducibility lies with the original authors to provide detailed instructions that
932 yield effective models.

933 *Appendix A.2.2. Template and Resolution Settings*

934 The highest resolution mesh settings available for each framework requiring
935 templates were employed. Specifically, for *DeepCSR*, which is not template-
936 based, a resolution setting of 256 was utilized for experiments and benchmark-
937 ing due to its stability. Although lower resolutions are available for each frame-
938 work and could significantly impact memory and GPU benchmarks, the focus

⁴<https://github.com/neuroneural/SurfaceReconstructionReview>

939 was on evaluating the performance of the highest-resolution meshes. Although
940 *DeepCSR* supports a resolution of 512, this setting was not explored further due
941 to excessive memory and GPU requirements that rendered it impractical within
942 the available computational environment.

943 *PialNN* utilized the white surface from *FreeSurfer* as the input template and
944 was constrained to *FreeSurfer*'s resolution. Frameworks such as *TopoFit*, *Corti-*
945 *cralFlow*, *CorticalFlow++*, *Vox2Cortex*, and *CortexODE* each required unique
946 templates corresponding to their respective resolutions, with the highest avail-
947 able resolution selected for each.

948 *Appendix A.2.3. ODE Solver Configurations*

949 For frameworks employing ODE solvers, the default number of steps and
950 solver types were utilized. Specifically, *CorticalFlow* and *CortexODE_euler* em-
951 ployed Euler solvers, while *CorticalFlow++* and *CortexODE_rk4* utilized Runge-
952 Kutta 4 (RK4) solvers. *CortexODE* was configured with a step size of 0.1 for
953 both solver variants, effectively demonstrating the strengths of the RK4 method
954 within their framework.

955 Both *CorticalFlow* and *CorticalFlow++* defined their step parameters in
956 terms of the number of steps, which were set to the default of 30 steps per
957 mesh. The Singularity training scripts were configured to reflect these settings.

958 *Appendix A.3. Hardware Configuration*

959 *Appendix A.3.1. Computational Benchmarks and Mesh Generation*

960 During computational benchmarks and mesh generation in evaluation mode,
961 the server's hardware configuration comprised an AMD EPYC 7551 32-Core
962 Processor, built on the x86_64 architecture. Using the Slurm job scheduler,
963 four cores were allocated for each framework's inference tasks per surface. The
964 processor operates at a clock speed of 2.532 GHz and features an extensive cache
965 system, including 2 MiB of L1d cache, 4 MiB of L1i cache, 32 MiB of L2 cache,
966 and 128 MiB of L3 cache.

967 The server is equipped with an NVIDIA GeForce RTX 2080 Ti GPU possessing
968 11,264 MiB of memory, operating with NVIDIA driver version 535.104.12
969 and CUDA version 12.2.

970 *Appendix A.3.2. Training Environment*

971 The training system is built on an x86_64 architecture running Ubuntu 20.04
972 OS. For each surface type and framework, four Intel Xeon Platinum 8468 CPU
973 cores with a base clock speed of 3.1 GHz were allocated during training sessions,
974 accommodating up to 16 cores for four surfaces. The CPU is equipped with a
975 4.5 MiB Level 1 data cache, a 3 MiB Level 1 instruction cache, a 192 MiB Level
976 2 cache, and a 210 MiB Level 3 cache.

977 One NVIDIA A100-SXM4-80GB GPU was allocated per framework, util-
978 izing NVIDIA driver version 535.104.12. The A100 GPUs require libraries
979 compatible with CUDA 11 or higher. The Slurm scheduler facilitated parallel
980 training of frameworks, dispatching each framework using Singularity containers
981 and scripts.

982 *Appendix A.4. Containerization and Code Management*

983 For security reasons, many high-performance computing centers prefer Sin-
984 gularity over Docker. Docker offers the capability to sandbox development and
985 production environments, providing unique sets of dependencies. Additionally,
986 Docker supports the deployment of microservices and facilitates the distribu-
987 tion of framework dependencies. Singularity provides similar functionalities to
988 Docker but operates with fewer privileges at runtime, offering enhanced security.

989 Each reviewed framework was forked, and corresponding Dockerfiles and
990 Singularity runtime scripts were created. Each framework repository includes
991 an updated README detailing the steps required to reproduce the results. The
992 companion repository for this study⁵ contains extensive documentation linking
993 to each fork, analysis scripts, and other pertinent information.

⁵<https://github.com/neuroneural/SurfaceReconstructionReview>